


20 Newsgroups, BERT-base, 8 epochs, 5 seeds (mean $\pm$ sd). Right: no corrupted labels at 0%.

Test accuracy, epoch picked on noisy held-out labels (%)



Corrupted labels predicted as given, last epoch (%)

